



Oss Ouahdi

Senior Product Marketing Manager assignment · Prepared for RevenueCat · September 2026

RevenueCat for Product Teams

How I refreshed the /for-product page: process, AI use and assumptions

Live page <https://revenuecat-for-product-nu.vercel.app/>
Hero variant B2 <https://revenuecat-for-product-nu.vercel.app/?v=b2>
GitHub repo <https://github.com/oss1010/revenuecat-for-product>
Walkthrough video [Link to come](#)

Contents

- | | | | |
|---|--|---|--------------------------------|
| 1 | Summary | 5 | Architecture and copy |
| 2 | Framing and research | 6 | Build and verification |
| 3 | Positioning, and every direction I killed | 7 | What I'd do next |
| 4 | The system: five agents and a context file | | Appendix: assumptions, sources |

1 Summary

The diagnosis. The current page was written for a RevenueCat that took monetization off a PM's plate: "You ship what matters, we handle monetization." The product outgrew that: Paywalls, Experiments with predicted 12-month LTV, web funnels and web billing hand monetization to PMs. The page, like competitors' comparison pages, still describes the earlier product. One comparison lists four shipped capabilities as missing.

The positioning. Monetization is a product decision, and RevenueCat lets product managers make it like one. The principle organizes the page, unstated. Three jobs carry it: change it without a release, learn what users value, fit the offer to every user, on a foundation that handles the rest. Revenue is the outcome, not the PM's job title.

The page. Hero: "Build a subscription experience worth renewing." Eight headlines carry the argument; five dynamic modules include an experiment simulator where conversion and predicted 12-month LTV pick different winners.

14

positioning directions
tested, 1 kept

19

hero lines considered,
1 kept

24

agent reviews run

56

claims failed by the audit
and fixed

2 **Framing and research:** RevenueCat's site, seven feature pages, five competitors

3 **Positioning:** 14 directions tested, 1 kept

4 **The system:** 5 agents, 1 context file, 1 pre-commit hook

5 **Architecture and copy:** 8 page sections, 5 dynamic modules, 19 hero lines

6 **Build and verification:** 8 phases, a red team, a final 48-claim sweep

How I worked

Tool	Role
Me	Direction and judgment: the diagnosis of the current page, the positioning, 13 directions killed and one kept, the copy, what the page had to prove and what it could leave out, and every final call. The AI widened the options and checked the work. The calls drew on my own product and developer marketing experience, including at Unity.
Claude, in chat	Strategy partner: positioning, copy, red-teaming its own advice, reviewing agent output, extracting RevenueCat animation frames for the builder's motion study
Claude Code	The builder: every file, the five agents, git, headless-browser verification
GitHub	Public repo, with the full decision history
Vercel	Hosting, redeployed every push
Chrome DevTools and Font Finder	Sampling RevenueCat's live typography and colors

Five review agents checked the work: pm-critic reads as three PMs, competitor-watch checks what competitors own, web-copy tightens copy, brand-guard checks RevenueCat's visual grammar, and claim-auditor traces every claim to a public source.

Three kinds of callout mark a collaboration running in every direction:

-  **I challenged the AI** where I overrode a recommendation
-  **The AI challenged me** where it pushed back and was right
-  **An agent caught it** where a review agent found what neither of us had

2 Framing and research

What I read first. RevenueCat's homepage, seven feature pages, pricing, changelog and docs (what shipped when), and the State of Subscription Apps 2026 report (plausibility checks). Then five competitors and two alternatives PMs weigh: building in-house, native store tools.

The competitive scan

	How it pitches product teams	Evidence tier	Signature claim	Scale
RevenueCat today	"You ship what matters, we handle monetization"	Owned: its own PM page, with a message the product has outgrown	Engineering time saved on in-app purchases	146K+ apps, \$17B+
Superwall	Dedicated PM page: own the monetization roadmap end to end, no app release or engineering queue	Owned (dedicated PM page)	Bills only on revenue through its own paywalls	\$1.6B+ revenue, 10,000+ apps
Adapty	Financial decisions on accurate revenue analytics. No PM page	Claimed (homepage section: testing for revenue, not just conversions)	12-month LTV and revenue prediction	Large, self-reported
Purchasely	The conversion problem is the whole first session, not one screen	Claimed	Onboarding journeys and personalized paywalls	Enterprise and media customers
Qonversion	A general PM page	Owned , weakly	No-code A/B testing	\$1B+ served revenue

	How it pitches product teams	Evidence tier	Signature claim	Scale
Apphud	Positioned as the RevenueCat alternative	Claimed	Revenue data accuracy	14,000+ apps, \$1B+ tracked, self-reported
In-house	No vendor fee, full control	n/a	"We already have engineers"	The default
Native store tools	Free and built in	n/a	Store-level experiments	Universal

What the scan showed. The category converged on one PM promise, now table stakes: change the payroll without engineering, test it, read the analytics. RevenueCat's edge is leadership, not a unique feature: the most complete platform, a 12-month forecast inside the experiment, unmatched scale.

The lesson: where a claim lives matters

Where the claim appears	What it means	Does it block RevenueCat from claiming it?
Dedicated persona page or homepage headline	Owned. They own that message to that audience	Yes. Don't lead with the same frame
Homepage section or feature page	Claimed. Strong but contestable	Only if RevenueCat can't prove it more completely
Docs, changelog, help center	Parity. A capability, not a position	No

I challenged the AI

When the competitor agent found revenue-based winners in two competitors' docs, the AI advised retreating to a claim nobody else could make: "grow revenue on the system that already runs it." I rejected it: documented capability is parity, not positioning, and leaders shouldn't position like challengers. We wrote the owned, claimed and parity rule into the agent.

3 Positioning, and every direction I killed

Positioning took five rounds, each killing something for a specific reason. The reasons are more useful than the directions.

Round	Directions tested	What killed them	Who killed them
✓ Kept	Monetization is a product decision. RevenueCat lets product managers make it like one	Structure survived all three agents; failed as copy, stays internal	Validated by agents
1	Monetization as a product surface; "speed is table stakes, knowing what worked isn't"; "your app store isn't your only store"	The AI recommended a hybrid: spine, proof, chapter. Three decisions wearing one coat	I did
2	"You own the number, now own the controls"; "change the payroll today, know what it's worth in 12 months"; "conversion is the wrong number"	First: near-verbatim on Superwall's PM page. Second: merged Superwall's speed and Adapty's forecast claims	The competitor scan
3	"Decide on money inside the experiment"; "grow revenue on the system that already runs it"	Two competitors' docs already had revenue winners; retreating cast a leader as a challenger	The agent, then me

Round	Directions tested	What killed them	Who killed them
4	"Control of every monetization decision, judged on long-term revenue"; "monetization calls a PM can defend"	First: Superwall's owned language; its reason to believe suited only existing customers. Second: defensive, from one synthetic persona's worry, not real PM pains	Three agents, then me
5	"Monetization that evolves alongside the product"; "make more money" and "long-term revenue"	First: a pace comparison and process, not an outcome. Second: the homepage's promise, not the persona's job	I did

✦ The AI challenged me

I wanted "autopilot": the system picking and shipping winners itself. The AI objected on three counts: rollout is human-triggered in the product, Autopilot is a competitor's product name, and it contradicts the control idea. Its reframe won: the machine recommends, the PM decides.

✓ An agent caught it

All three PM personas rejected "monetization is a product decision" as copy: a truism, a turf claim, wrong for enterprise. Yet each, unprompted, described one of our three jobs in their own words. Structure validated, phrase killed: unowned is not the same as resonant.

An honest note. The final idea is close to round one's. What changed is underneath: it now comes from the PM's job, is structured by their real pains, and is differentiated by proof, not a clever phrase.

From seven pain points to three jobs

Seven PM pains map onto three jobs and a foundation, each a page section.

The PM's pain	Job	What delivers it	Page section
Paywall changes queue behind engineering	Change it without a release	Paywalls, remote configuration, one-action rollout	4, 5
Not knowing which test result to trust	Learn what users value	Experiments with predicted 12-month LTV	5
Blind spots between acquisition and revenue	Learn what users value	Web funnel campaign tracking; Apple Search Ads keyword filters in Charts	5, 6
Monetization problems spotted too late	Learn what users value	Paywall charts, realized LTV per customer, Benchmarks	5
One offer for every user	Fit the offer to every user	Targeting, funnels that branch	6
Store commission pressure	Fit the offer to every user	Web billing	6
Store policy changes and refund abuse	Foundation	Refund Control, Apple and Google API changes handled	7

4 The system: five agents and a context file

I built a small review system, not a prompt chain, so every claim, headline and visual was checked by something other than its author.

Agent	Its job	What it caught
pm-critic	Reads as three PM archetypes	Validated the three-job structure; raised the enterprise governance question and product gaps
competitor-watch	Checks what competitors own	Killed a Superwall headline I planted as an unannounced calibration test. Flagged v0.1's core idea as Superwall's language
web-copy	Drafts and tightens copy	All 14 proposed cuts applied pre-build: about 85 words (11%)
brand-guard	Checks RevenueCat's visual grammar	Five phase 3B fixes, all applied: a leftover CSS rule breaking the funnel layout; page section 3 tabs, the largest blue fill; page section 7 icon tiles reading as the old feature grid; testimonials reworked into one centered block; a container narrowed from 1,200px to about 990px, RevenueCat's width. Verdict: reads as a RevenueCat page
claim-auditor	Traces claims to public sources	56 fails in total, 4 in the final 48-string sweep. Among them: a nonpublic product name, an unattributed testimonial, "one click" where the product says "one action," Figma designs that arrive via a separate plugin, a Benchmarks band at 25th to 75th percentile where RevenueCat uses 20th to 80th, and "win-back," which RevenueCat uses for users already gone

The context file. CLAUDE.md holds audience, non-negotiables, design tokens, contrast rules and color roles. Every session and agent reads it first.

Context states intent. Tooling enforces it. CLAUDE.md was created with an em dash, breaking its own rule, so the rule became a pre-commit hook. The AI in chat logged the hook as built before it existed, then caught it. Now it's real, and it blocks.

I challenged the AI

What the agents missed: a page for product managers that never said "product manager" in visible copy. Four agents reviewed it; none was asked who it speaks to. I caught it reading through. Agents check what they're told to check.

5 Architecture and copy

Architecture decisions

- **The hero carries the idea; the sections carry the proof.**
- **Half static, half dynamic.** The current page is all static imagery, and visual storytelling is judged first.
- **One fictional app across every visual,** Tidelark, a premium AI morning-routine coach: one story, not four unrelated ones.
- **Zigzag layout** for the job sections, like RevenueCat's feature pages.
- **"Learn more" links to matching RevenueCat feature pages,** so PMs can go deeper.

The hero: 19 lines considered, 10 shown

Candidate	Verdict
"Build a subscription experience worth renewing."	 Kept. A subscription is renewal by definition, so "worth renewing" says user value in two words. Top combined persona score
"Build the reason they renew."	Cheekiest line; lands best once the argument is made. Became the closing line and hero test variant
"Your paywall shouldn't wait for a release."	Strongest felt pain, but table stakes and close to Superwall's promise. Became page section 4's headline

Candidate	Verdict
"Change the paywall today. See its 12-month value."	Two unlinked ideas, and "change" is an activity, not a capability
"Test pricing like you test features."	No outcome, no benefit, no difference
"Ship the paywall that pays, not the one that converts."	A competitor claims revenue over conversion; "pays" reads as a guarantee
"Own your paywall without waiting on engineering."	Superwall's positioning, planted to test competitor-watch. It caught it
"Launch a paywall today. Judge it on the next 12 months."	"Today" isn't true for a first paywall; the first half sits on Superwall's promise
"Stop shipping paywalls that win the week and lose the year."	Sharpest idea, but opens by telling readers they're wrong. Became page section 5: "Win the year, not just the week."
"Build the subscription experience your users stay for."	The right idea. "Paywalls" narrowed it, "subscription experience" widened it

I challenged the AI

The AI's first three hero lines were weak, and I said so. Short, accurate and unowned, they met the rules but didn't lead with what the reader gets. The winning wordplay was mine.

The skimmer test

The headlines alone should carry the argument; the list doubles as the page map.

Page section 1	Build a subscription experience worth renewing.
Page section 2	Trusted by teams at [logos]
Page section 3	Change faster. Learn what renews. Fit every user.
Page section 4	Your paywall shouldn't wait for a release.
Page section 5	Win the year, not just the week. The test ends. The learning doesn't.
Page section 6	Different users. Different offers. App or web.
Page section 7	Subscribe once. Unlocked everywhere.
Page section 8	Two ways to start.
Closing line	Build the reason they renew.

The hero animation

The rebuilt loop follows the hero's sub in about eight seconds: **Design** changes the featured plan, **Test** runs weekly against monthly with B pulling ahead, **Roll out** fills a bar to 100%: "No app release."

I challenged the AI

Version one's cursor clicked three cards that changed nothing: decoration, not a demo. I also pushed for a paywall like a real, premium app's, not two plan rows. Now every click visibly changes what the user sees.

The three-jobs module, three versions

First, three arrow rows: a table of contents posing as a section. Then tabs with a before-and-after slider, cut for a weak reveal. Then side-by-side panels. **Change**: before, Engineering and the App Store own most steps; with RevenueCat, every step says "You" and reaches "Live" first. **Learn**: paying customers as dots, week one favoring A,

the predicted year B. **Fit:** one payroll for everyone against the right offer per audience, app and web.


I asked for B's dots to stay lit. The builder kept them fading, far less than A's, as the model says. Honest proportions beat a cleaner picture.

6 Build and verification

Stack. Static HTML, CSS and vanilla JavaScript on Vercel. One landing page needs no framework, and fewer moving parts meant fewer ways to break the night before the deadline.

Phase	What got built	Checkpoint
1	Scaffold, locked copy, design tokens, first simulator	Worked first time
1.5	Visual grammar, proof points, simulator v2	Colors sampled from RevenueCat's pages
2	PM copy, three-jobs tabs, first hero loop, zigzag layout	Modules verified by keyboard
3A	Shared paywall, simulator v3 with forecast chart, new hero	Every illustrative number refit to published ranges
3B	Charts module, visuals for page sections 3 to 7, brand-guard review	No placeholders left
4	Tighter hero, side-by-side three-jobs panels, compact simulator, line charts, web billing	Mobile simulator nearly halved
5	Red team; page section 7 rebuilt as the foundation; final three-agent review	48 claims swept, 4 fixed
5.1	Final fixes from the agents' ranked list, including moving the Pixelcut quote	Design frozen

The simulator, four versions



The experiment simulator: judge by conversion or by predicted 12-month LTV, and the winner flips.

Version	What it did	What was wrong	Who caught it
v1	A toggle swapped one metric and the winner flipped	Tension spanned only two moments. The rollout message sat on the user's phone, where users would never see it	The AI's review
v2	Both metrics per card; the PM's panel split from the user's phone; a real paywall; a one-time cursor demo	It showed what won, not why. Numbers too small to feel real; the decision too obvious	I did
v3	Forecast chart, observed month solid, predicted months dashed; B overtakes at month 4. A "why B wins" panel. Live traffic split on the phone	LTV averaged across every test customer read like a monthly figure; the module topped 2,100px on mobile	I did
v4	LTV per paying customer (\$59.52 against \$130.46), math under the chart; the "why" on the chart lines; two-step rollout; cursor demo runs the whole sequence	Final. Mobile height down from 2,167px to 1,210px	

✦ The AI challenged me

My v3 idea was a crossover: A leads three months, then B overtakes. Impossible with annual versus monthly: annual plans are paid upfront, so B leads from day one, as anyone who knows subscription data would spot. Switching to weekly versus monthly, one of subscription apps' most debated pricing questions, made the story true.

✓ An agent caught it

Even then, illustrative numbers must be plausible. The audit refit every input to RevenueCat's published conversion and retention ranges, and the story held: A leads early, B overtakes at month 4 and ends about 47% higher.

✦ The AI challenged me

I wanted the winner to roll out automatically once the forecast appeared, sparing visitors a click. But that's autopilot, which we had rejected. So the cursor demo performs the rollout: every visitor sees it, and the PM still decides.

Brand: rebuilt, not reused

I studied seven RevenueCat feature pages and two product animations, then rebuilt the visual grammar in code, not their files: the work had to be mine, and RevenueCat would recognize its own assets. Colors and logos came from the press kit, checked against the live site. In production, I'd use the design library and brief the agency for new work.

- Surface colors sampled pixel by pixel from RevenueCat's pages: off-white #F9F9FB, a lavender and peach hero wash
- Color roles: blue for actions, green for winning and live states, red for emphasis
- RevenueCat's fonts are commercially licensed, so Manrope sets display and RevenueCat's declared fallback stack sets body

✓ An agent caught it

Three accessibility failures in the published brand palette: white on brand red, brand green on white, brand blue as link text on off-white. Buttons went blue, as on the live site. Green became a fill, not text. Links got a darker blue that passes everywhere.

Verified: 390px first, then 1440px, in a headless browser when the preview wasn't visible. Keyboard, screen reader announcements, focus handling, reduced motion, every outbound link. Fixed: both paywall states showing at once, compound words breaking at hyphens on mobile, a toast covering the status line.

The red team: the new page against the current one

A red team attacks your own work the way a skeptical reviewer would. Here, that was me and Claude first, then all three review agents.

	The current page	The new page
Promise	"You ship what matters, we handle monetization." Hand it off	"Build a subscription experience worth renewing." Own it
The PM's role	Bystander: RevenueCat saves engineering time	Operator: change it, learn from it, fit it
Structure	Seven feature tiles, then SDK setup steps	Three jobs from the PM's pains; the headlines alone tell the story
Proof	One CTO quote about engineering time	Pixelcut, Floga, OpenAI, and RevenueCat's scale: 146K+ apps, \$17B+ processed
Visuals	Static screenshots	Hero loop, a simulator where the winner flips, live charts, web-to-app journey
New capabilities	One A/B testing tile	AI Editor, predicted 12-month LTV, one-action rollout, Funnels, web billing, Benchmarks

What it found, and what changed:

- **Page section 7 broke the story.** "Decide who can change a live payroll" was a governance feature repeating a page section 4 bullet beside an unrelated visual. It became the foundation: "Subscribe once. Unlocked everywhere."
- **Two truths from the current page had gone missing:** Customer Center, and Apple and Google API changes handled for you. On a page promising "worth renewing," losing the retention feature was a real gap. Both returned in page section 7.
- **One pain had no home.** Blind spots between acquisition and revenue now map to web funnel campaign tracking.
- **Density.** The web-to-app journey went from seven steps to four at half the height; tabs stopped scrolling sideways on mobile.

All three agents then reviewed the finished page. pm-critic: a clear win over the current page for solo and growth PMs, partial for enterprise, prompting a security and compliance link on the Enterprise card. brand-guard: no contrast or color-rule failures. claim-auditor: 48 strings, 4 failed, all fixed.

✓ An agent caught it

Pixelcut's quote, beside the simulator, celebrated a 16% increase in subscribers; its case study credits a weekly plan. So it praised the early win the simulator warns can mislead, as anyone at RevenueCat who knows the case would spot. Moved to page section 4, it proves fast iteration instead of undercutting the forecast.

7 What I'd do next

With more time

- **Put it in front of real PMs.** Five PMs from target accounts and existing customers, one task: read only the headlines, then say what RevenueCat does for them. Their words, not synthetic personas, would set the next round of copy.
- **Run the first two tests.** The hero, 'worth renewing' against 'the reason they renew', on sign-up rate from the page. And the link for existing customers, on dashboard clicks, then first payroll published.

- **Build a guided demo.** The simulator is a small version of it: pick your app category, build a paywall from a template, run a simulated test, roll out the winner. Five minutes, no account, no SDK. The current page's demo dashboard link now lands on a login screen; this would replace it. One rule: it answers only from RevenueCat's own data and docs.

Three questions I'd bring to Product

All three are about trust in the forecast, the page's central claim.

- Could there be a public page on how the 12-month forecast is made? The audit found none, and it's the first page I'd write.
- Could the forecast show a confidence range? Credible intervals currently cover conversion metrics only.
- Could benchmarks sit inside the decision? Showing how a variant compares with similar apps answers 'is my number any good?', a question only RevenueCat's scale can answer.

A Appendix

Assumptions

- 1 The page serves prospects and existing customers in an unknown mix. It leads for prospects, with one link for customers.
- 2 Primary conversion: self-serve sign-up. Sales is secondary.
- 3 PMs own the product experience and strategy. Revenue-number ownership varies by company size.
- 4 Only public, shipped capabilities are claimed, none from a roadmap.
- 5 All product data is illustrative and labeled. Tidelark is fictional.
- 6 RevenueCat's fonts are licensed; the page substitutes.

What I couldn't verify

Real traffic, customer interviews and internal data. The PM personas are synthetic, their feedback treated as hypotheses.

Key sources

RevenueCat pages

revenuecat.com/
revenuecat.com/for-product/
revenuecat.com/feature/experiments

Docs and changelog

revenuecat.com/changelog/release/see-predicted-12-month-ltv-winners-in-experiments-2026-02-12
revenuecat.com/changelog/release/roll-out-an-experiment-winner-in-one-action-2026-02-06

Every source the audits used is in `docs/iteration-log.md`.

revenuecat.com/changelog/release/compare-subscription-metrics-against-industry-benchmarks-2026-04-13

Competitors

superwall.com/solutions/product-managers
adapty.io/
apphud.com/revenuecat

The report

revenuecat.com/state-of-subscription-apps